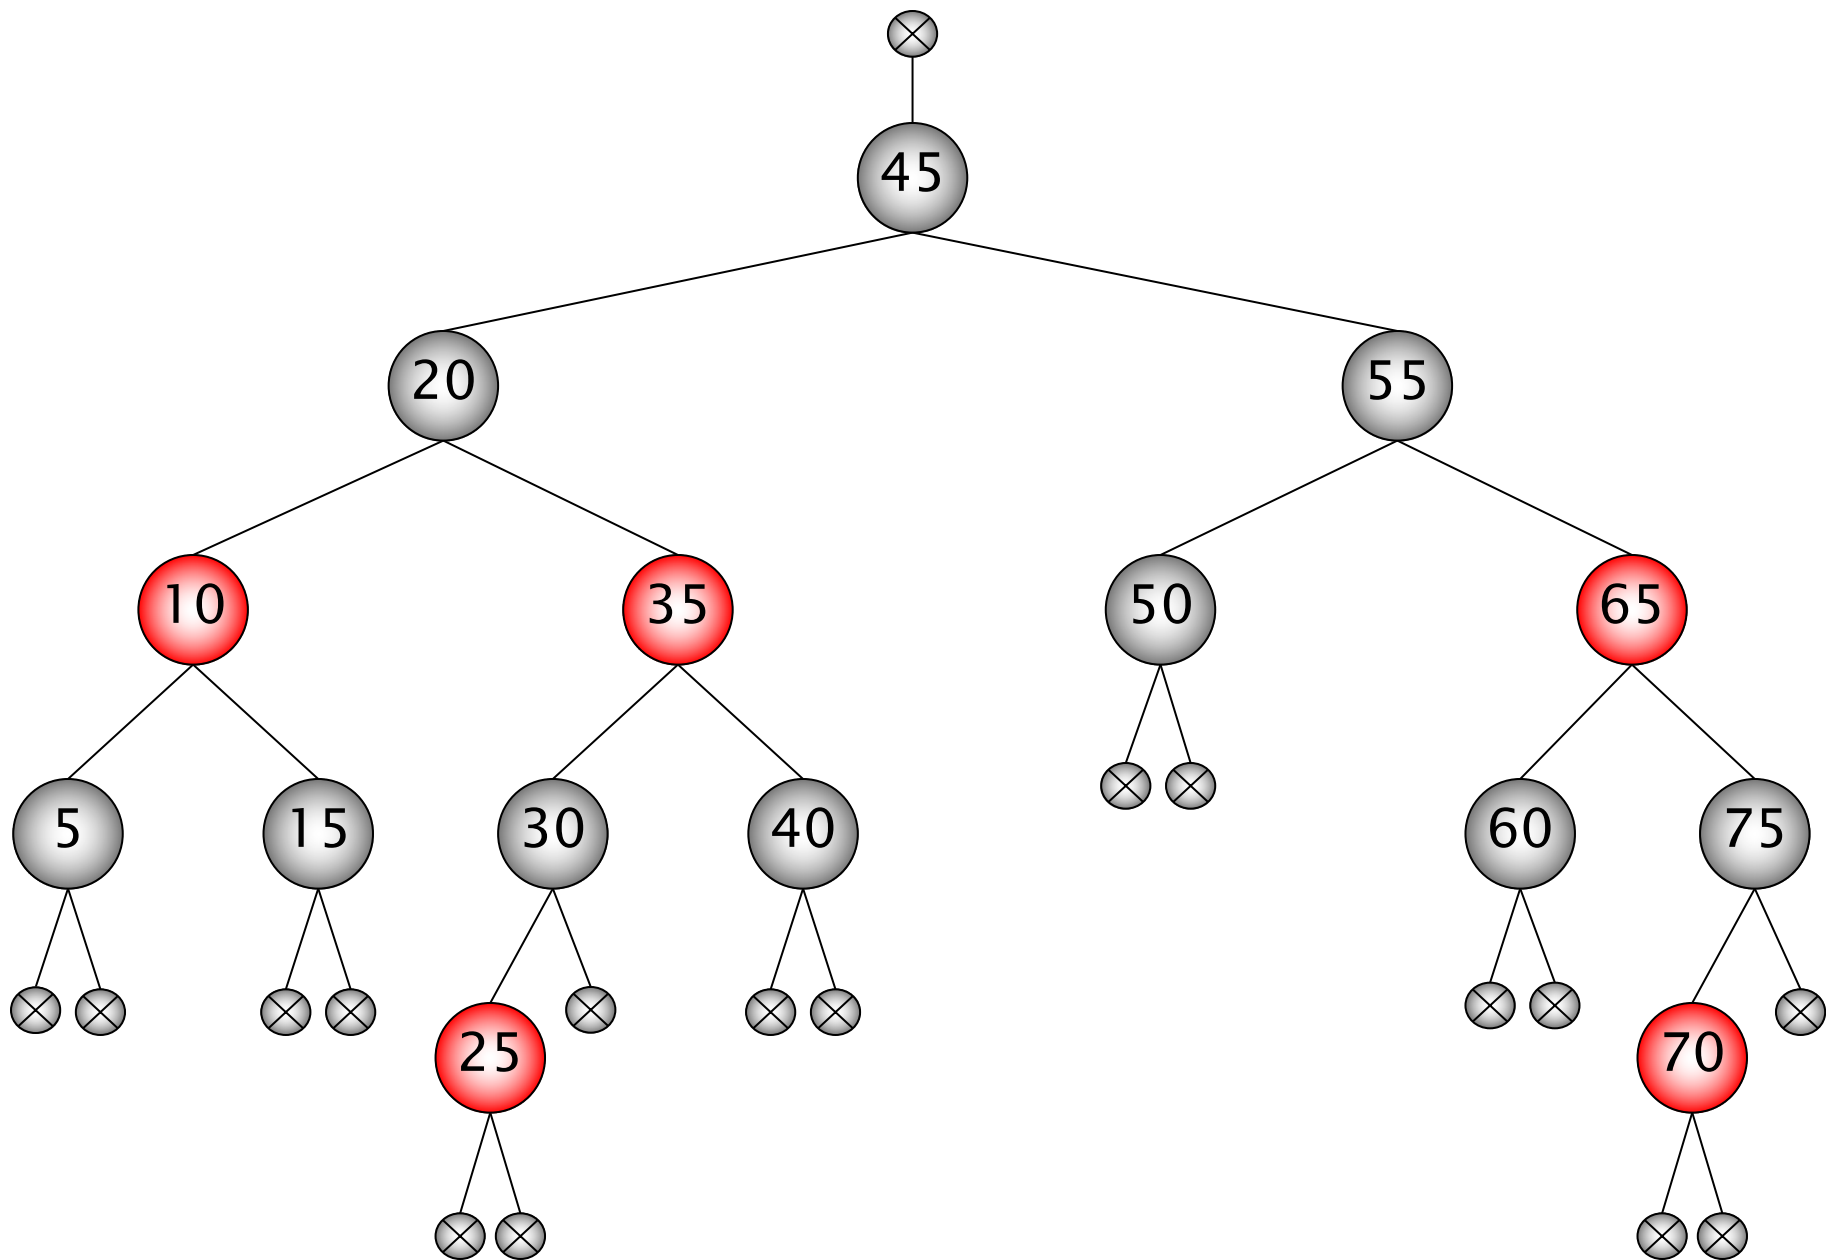
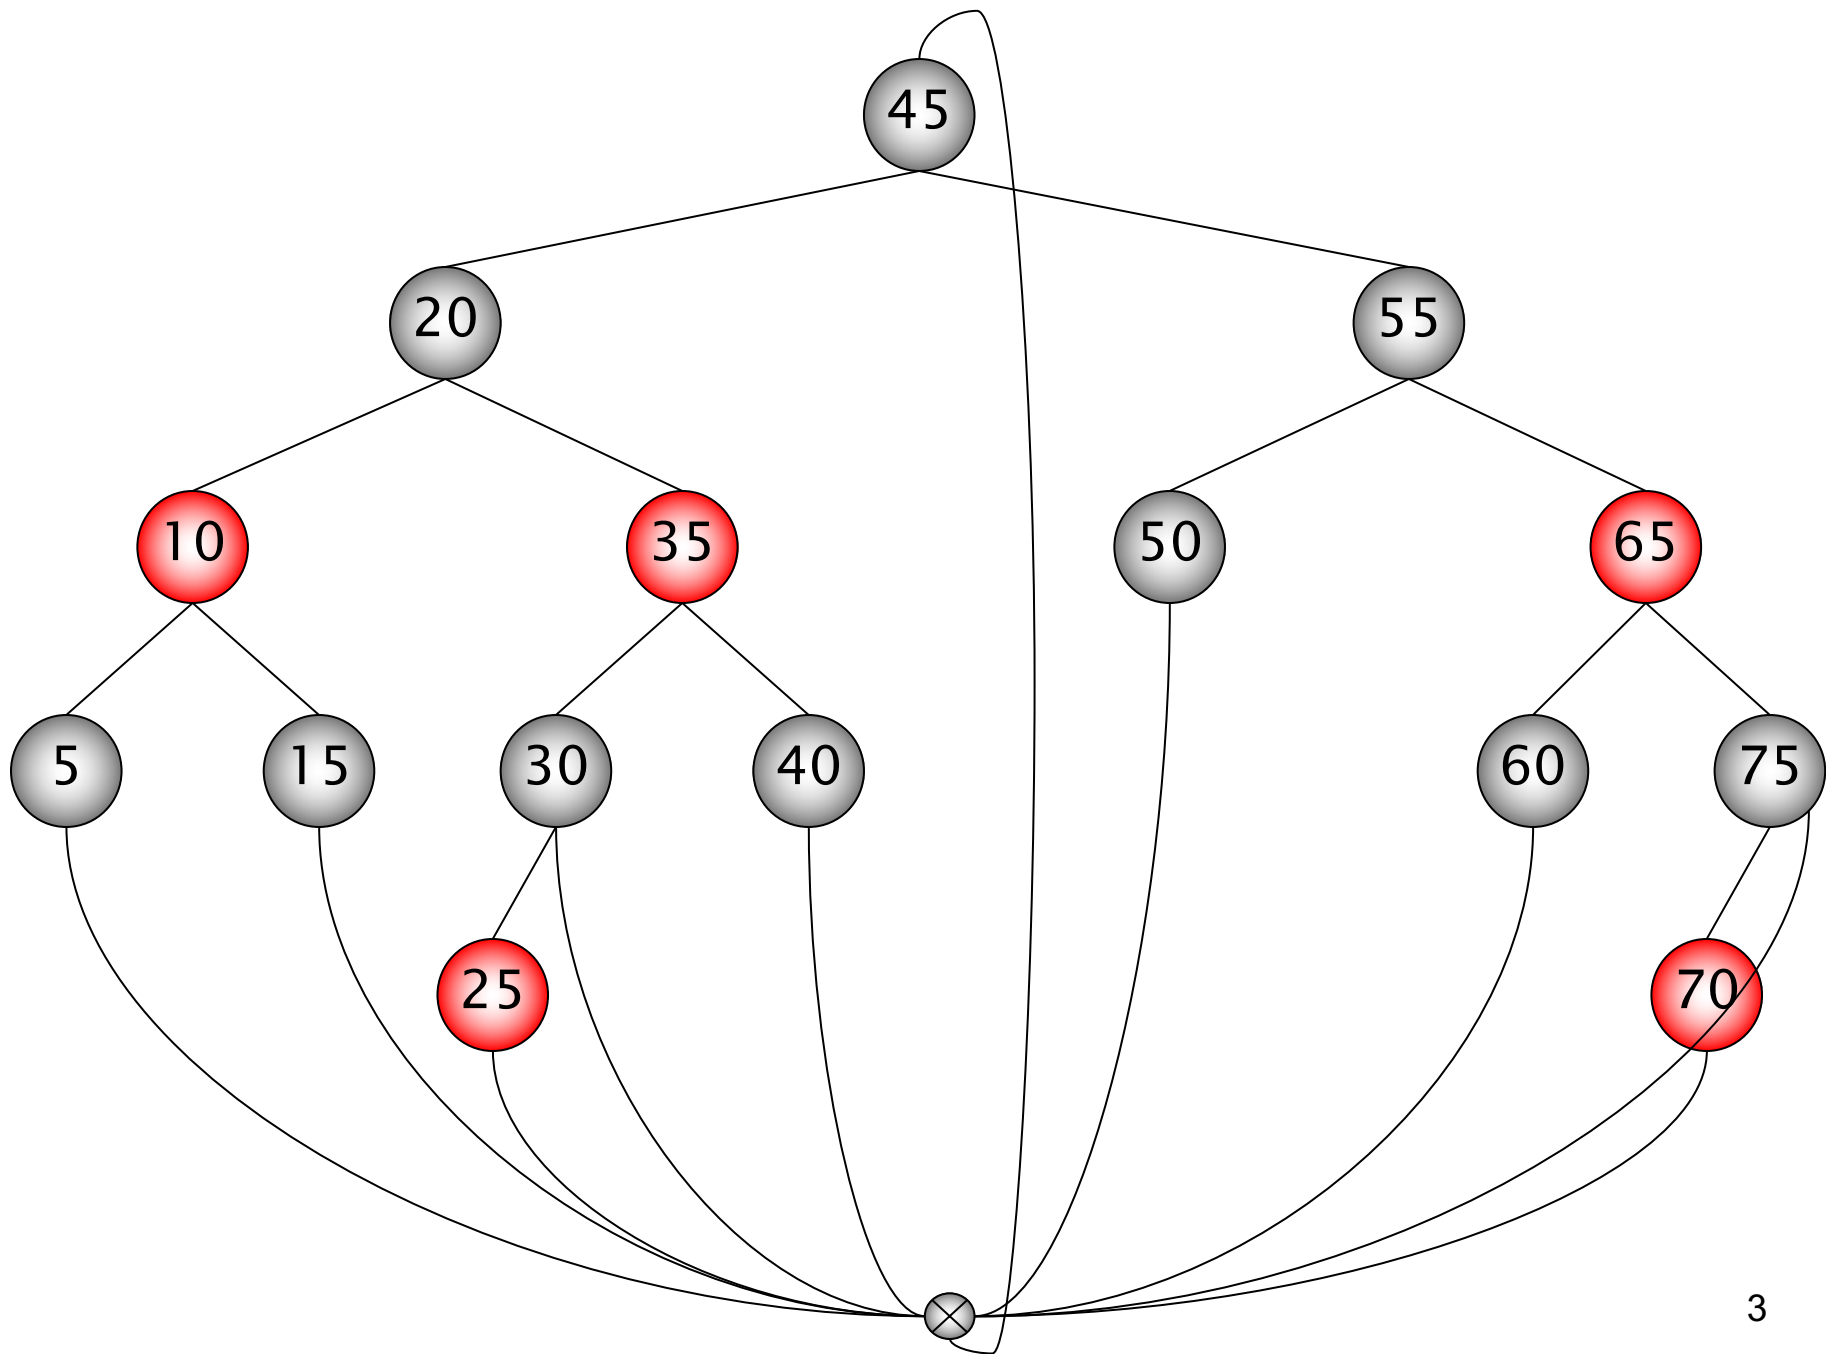


Red-Black Trees

- A *red-black tree* (RBT) is a type of self-balancing BST
 - Each node in an RBT is labeled as *red* or *black*
- Operations on RBTs take $O(\log n)$ time in the worst case since the height of an RBT is at most $2 \log_2(n + 1)$
 - The height of an AVL tree is at most $1.44 \log_2 n$
- However, a careful nonrecursive implementation can be done relatively effortlessly compared with AVL trees





Definition of Red-Black Trees

An RBT is a BST that satisfies the following *red-black criteria*:

1. Every node is either *red* or *black*
2. The root of the tree is always *black*
3. All leaves are *black*
4. If a node is *red*, then its parent is *black*
5. Any path from a node to any of its leaves contains the same amount of black nodes, called *black height*

Representation of Red-Black Trees

```
typedef struct Node * Ref;
struct Node {
    int key;
    int color;
    Ref parent, left, right;
};

ref getNode(int key, int color, Ref nil) {
    p = new Node;
    p->key    = key;
    p->color  = color;
    p->left   = p->right = p->parent = nil;
    return  p;
}
```

The Initial State of a Red-Black Tree

```
Ref nil, root;
```

```
...
```

```
nil = new Node;
```

```
nil->color = BLACK;
```

```
nil->key = -1;
```


```
nil->left =
```

```
nil->right =
```


```
nil->parent = nil;
```

```
...
```

```
root = nil;
```



Tree Rotation



```
leftRotate(Ref & root, Ref x) {
```

```
    y = x->right;
```

```
    x->right = y->left;
```

```
    if (y->left != nil)
```

```
        y->left->parent = x;
```

```
    y->parent = x->parent;
```

```
    if (x->parent == nil) root = y;
```

```
    else
```

```
        if (x == x->parent->left)
```

```
            x->parent->left = y;
```



```
        else
```

```
            x->parent->right = y;
```

```
    y->left = x;
```

```
    x->parent = y;
```

```
}
```



Search Operation

- An RBT is a BST so the search algorithm for an RBT is the same as the search algorithm for a BST

Insertion Operation

- The new node, as usual, is placed as a leaf in the tree
- The node must be colored *red*
 - If the parent is *black*: The red-black criteria are maintained
 - If the parent is *red*: The operation causes a violation of the criterion “no consecutive red nodes”
 - ➡ Rebalancing the tree is needed

Insertion Operation: Pseudo-code


```
void    RBT_Insertion(Ref & root, int key) {  
    x = getNode(key, RED, nil);  
    BST_Insert(root, x);  
    Insertion_FixUp(root, x);  
}
```

```

BST_Insert(Ref & root, ref x) {
    y = nil;
    z = root;
    while (z != nil) {
        y = z;
        if (x->key < z->key)          z = z->left;
        else if (x->key > z->key)      z = z->right;
        else                          return;
    }
    x->parent = y;
    if (y == nil)    root = x;
    else
        if (x->key < y->key)  y->left = x;
        else                y->right = x;
}


```

Some Conventions




- x : The pointer designated to point to the newly added leaf
- p : The pointer designated to point to the *parent* of the node pointed to by x
- u : The pointer designated to point to the *uncle* of the node pointed to by x
- g : The pointer designated to point to the *grandparent* of the node pointed to by x


Imbalance Cases



Case 1



Case 2




Case 3

The Strategies for Rebalancing an RBT

- Case 1


→ Reverse the color of three nodes: u , p , and g



- Case 2

→ Reverse the color of two nodes: p and g


→ Run a rotation at g




- Case 3

→ Run a rotation at parent p


→ Go to Case 2



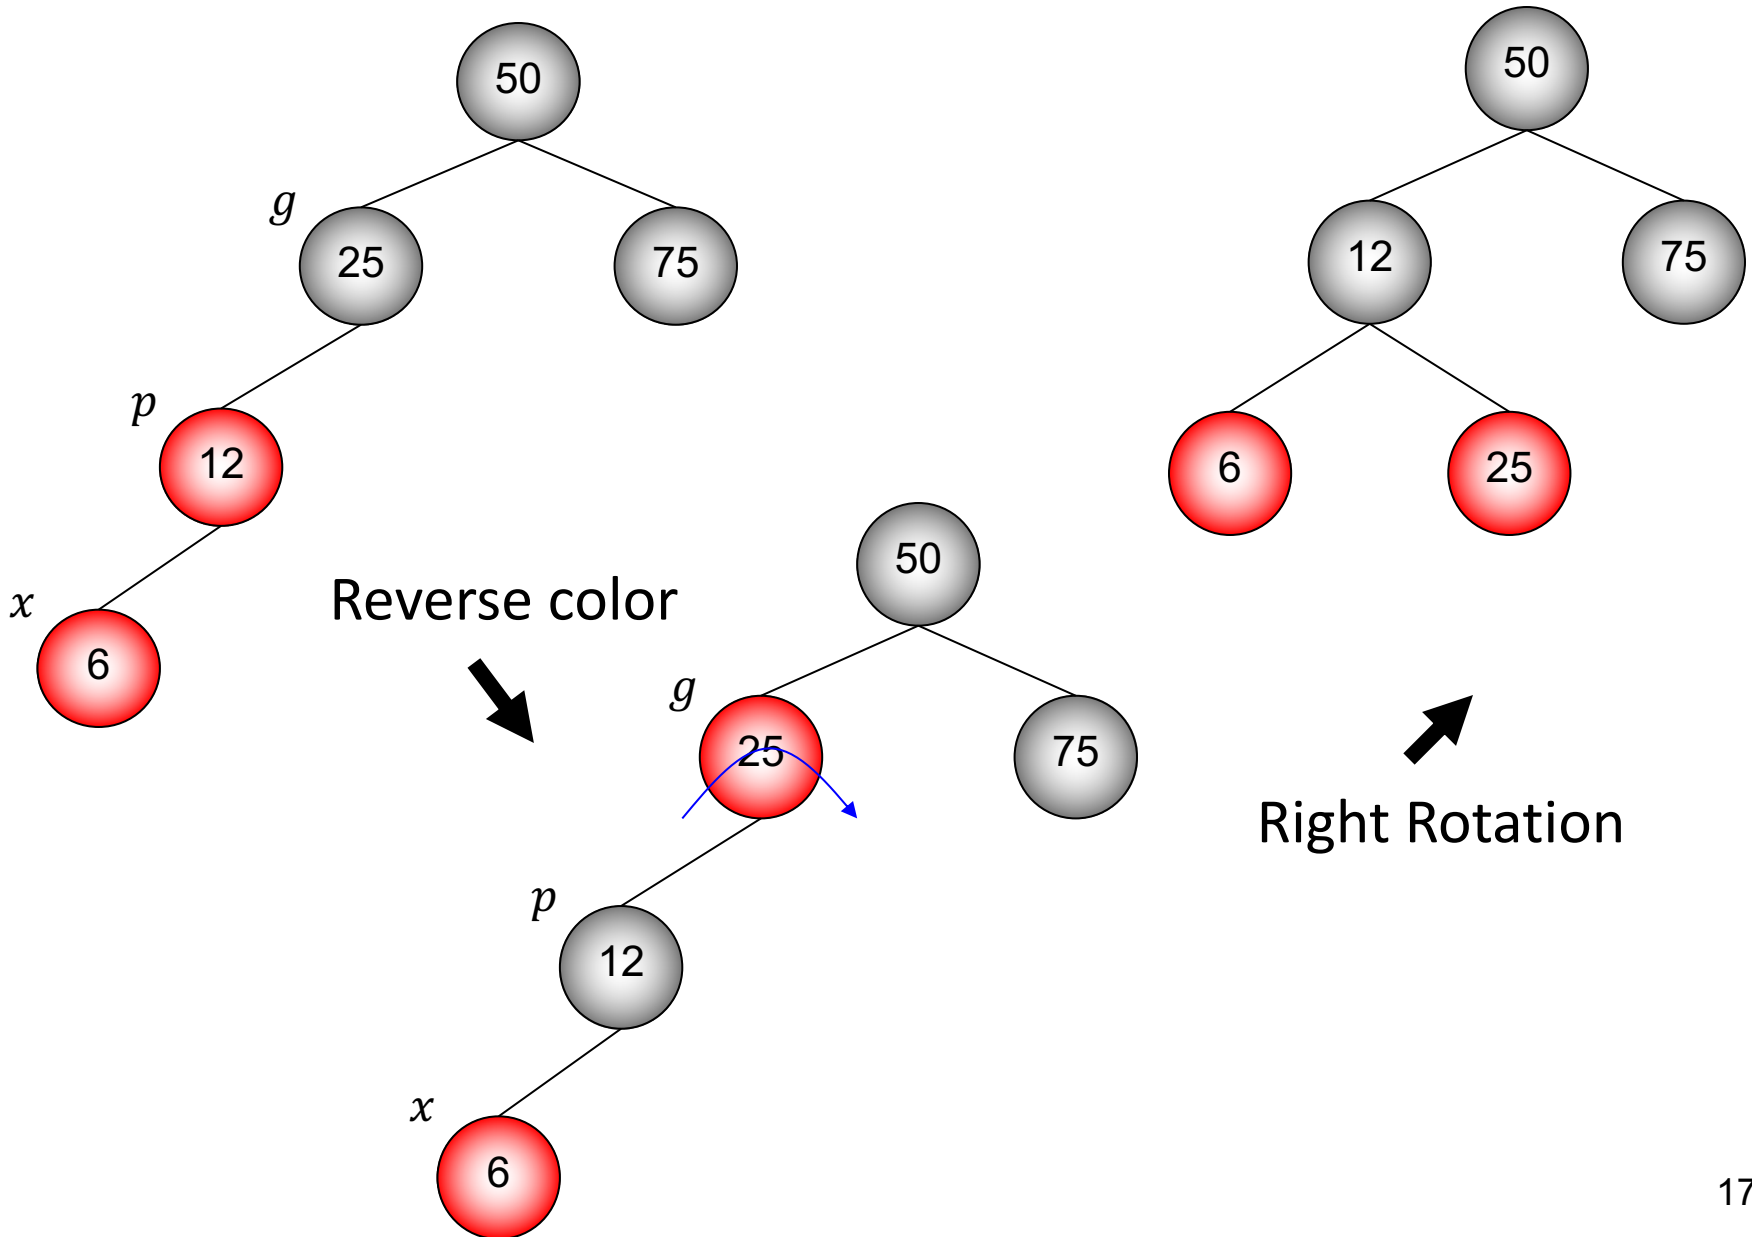
Case 1




Reverse color




Case 2




Case 2




Case 3





Case 3



Right Rotation



Reverse color



Left Rotation

Insertion Operation: Pseudo-code

```
Insertion_FixUp(Ref & root, Ref x) {  
    while (x->parent->color == RED)  
        if (x->parent == x->parent->parent->left)  
            ins_leftAdjust(root, x);  
        else  
            ins_rightAdjust(root, x);  
    root->color = BLACK;  
}
```

```

ins_leftAdjust(Ref & root, Ref & x) {
    u = x->parent->parent->right;
    if (u->color == RED) {
        x->parent->color = u->color = BLACK;
        x->parent->parent->color = RED;
        x = x->parent->parent;
    }
    else {
        if (x == x->parent->right) {
            x = x->parent; leftRotate(root, x);
        }
        x->parent->color = BLACK;
        x->parent->parent->color = RED;
        rightRotate(root, x->parent->parent);
    }
}

```

Deletion Operation


- Let r be the node to be deleted. One of the following three cases will arise:
 1. r is a leaf
 2. r has one nonempty subtree
 3. r has both nonempty subtrees
- Let's consider the third case:
 - Find the predecessor or successor, say s , of r
 - Copy the data of s into r *except the color*
 - ➡ Now, the node to be deleted is s

```

RBT_Deletion(Ref & root, int k) {
    z = searchTree(root, k);
    if (z == nil) return;
    y = (z->left == nil) || (z->right == nil) ?
        z : Predecessor(root, z);
    x = (y->left == nil) ? y->right : y->left;
    x->parent = y->parent;
    if (y->parent == nil)    root = x;
    else
        if (y == y->parent->left) y->parent->left = x;
        else                    y->parent->right = x;
    if (y != z)    z->key = y->key;
    // if (y->color == BLACK) Del_FixUp(root, x);
    delete y;
}

```


Some Conventions



- y : The pointer designated to point to the node to be deleted
- x : The pointer designated to point to the child of the node pointed to by y
- p : The pointer designated to point to the parent of the node pointed to by y
- w : The pointer designated to point to the sibling of the node pointed to by y

What'll Happen After Deleting a Node?

- If y was red, the red-black criteria still hold because:
 - No black heights in the tree have changed
 - No red nodes have been made adjacent
 - If y was black:
 - The number of black nodes on every path passing through node y is decreased by 1
 - The “no two consecutive red nodes” criterion is also violated if both nodes, p and x , are red
- ➡ Rebalancing the tree is needed

```

RBT_Deletion(Ref & root, int k) {
    z = searchTree(root, k);
    if (z == nil) return;
    y = (z->left == nil) || (z->right == nil) ?
        z : Predecessor(root, z);
    x = (y->left == nil) ? y->right : y->left;
    x->parent = y->parent;
    if (y->parent == nil)    root = x;
    else
        if (y == y->parent->left) y->parent->left = x;
        else                    y->parent->right = x;
    if (y != z)    z->key = y->key;
    if (y->color == BLACK) Del_FixUp(root, x);
    delete y;
}

```

Black Token

- *Black token* is an abstract concept
 - It's used to explain the meaning of the rebalancing process
- If y is black then x will receive black token
 - If x is black then x is called *doubly-black node*
 - If x is red then x is called *red-black node*
- What is the role of black token?
 - After deleting y , all black heights passing through x are one unit shorter than the others
 - The occurrence of black token logically implies that every black height passing through x is supplemented by one black

Black Token

- The black token tends *to be pushed toward the root* of the tree
- It could (i) be neutralized on the way to the root or (ii) finally be attached to the root, thus making the root as a doubly-black node



In both cases, the rebalancing process stops immediately

The Strategies for Rebalancing an RBT

- There are 4 cases that may happen when a black node is deleted
- We focus our attention on node x , the one with black token

Case 1



- If x is a red-black node
 - The node will be neutralized by coloring it black
- If x is the root
 - The black token will be eliminated

Case 2

- Node x is doubly-black
- Its sibling w is black
- Both of w 's children are black

Solution:

- Reverse the color of sibling w
- Attach black token to parent p

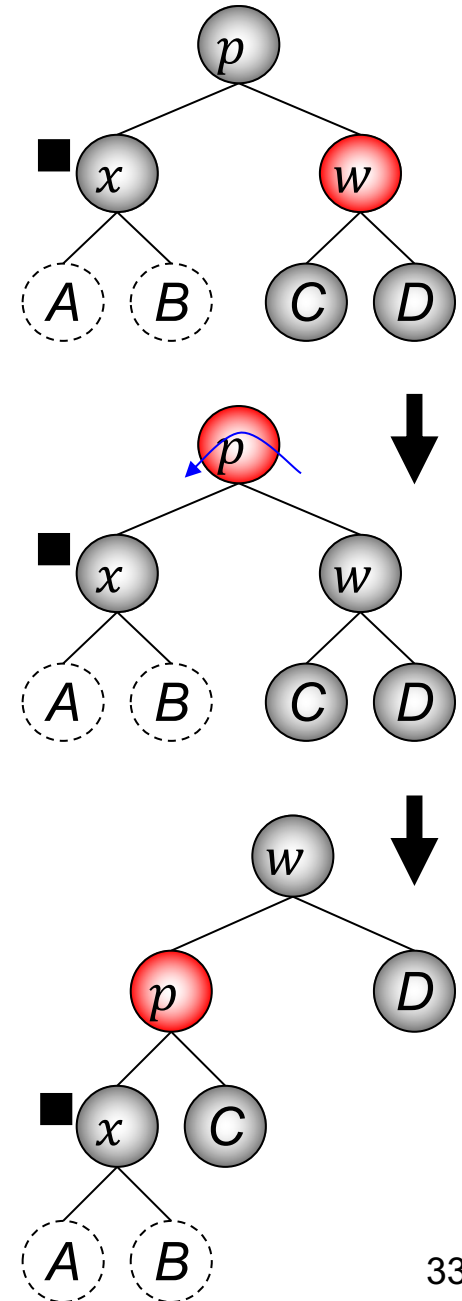


Case 3

- Node x is doubly-black
- Its sibling w is red


Solution:

- Reverse the color of parent p and sibling w
- Perform a rotation at parent p



Case 4

- Node x is doubly-black
- Its sibling w is black
- At least one of w 's children is red



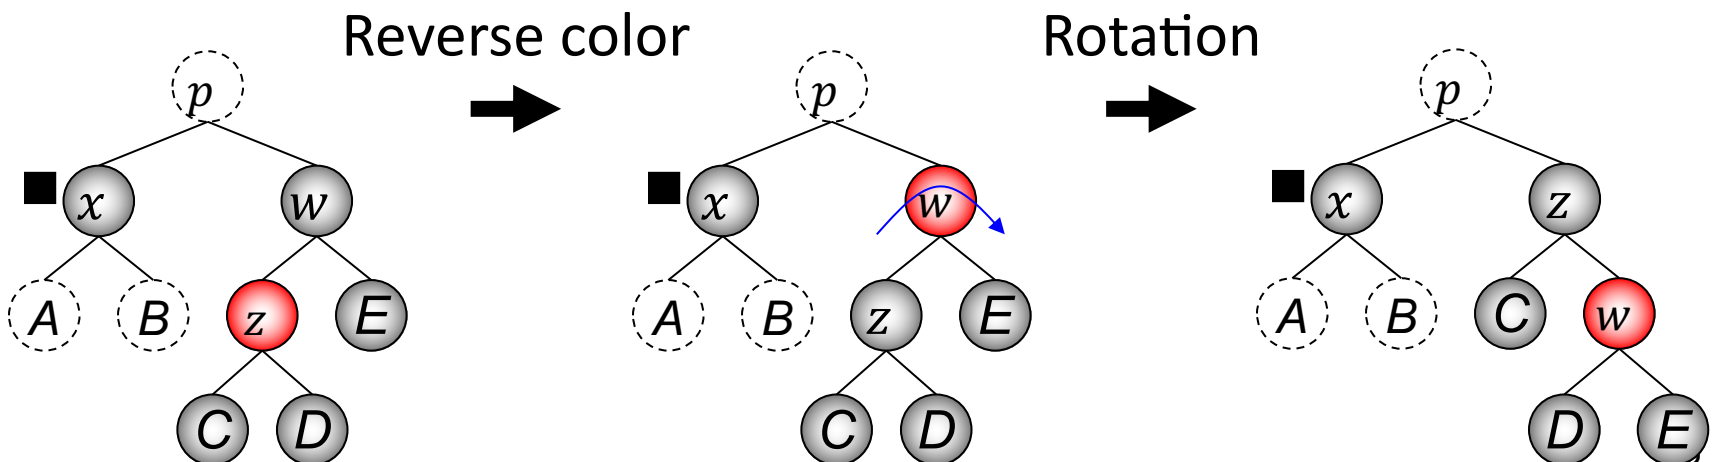
➡ The solution depends on the w 's child node which is the external grandchild of grandparent p

Case 4.a

- The external grandchild is black

Solution:

- Reverse the color of internal grandchild z and sibling w
- Perform a rotation at the sibling w
- Go to Case 4.b

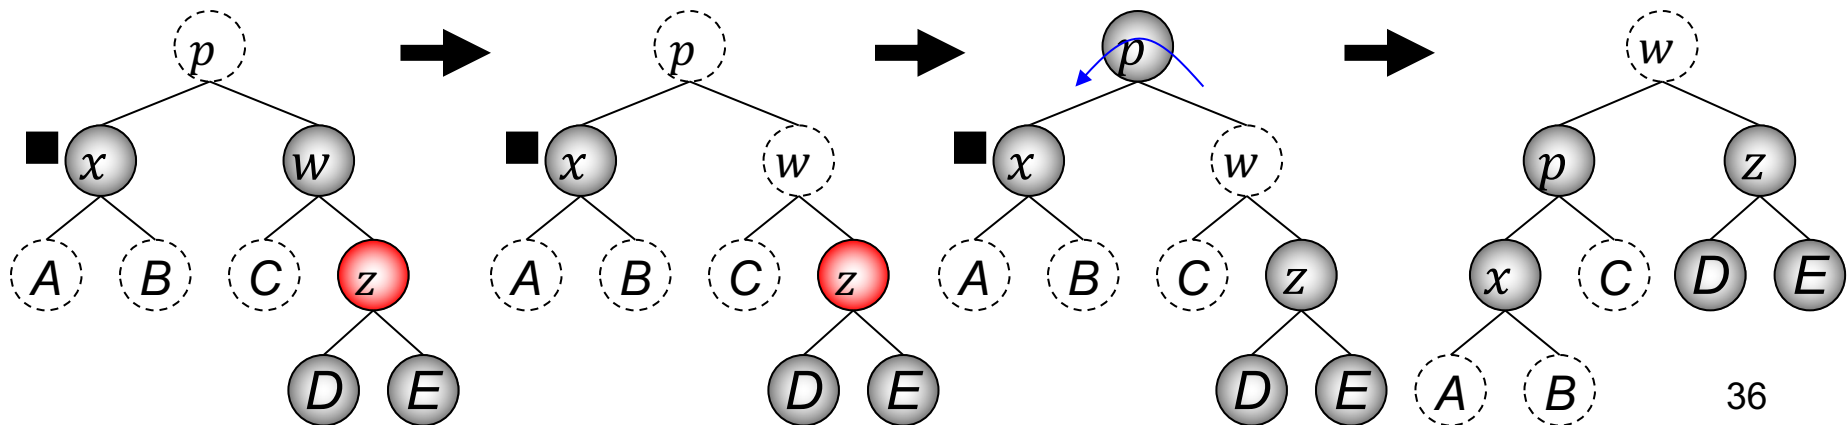


Case 4.b

- The external grandchild is red

Solution:

- The sibling w inherits the color of its parent p
- The color of grandparent p and external grandchild z is set to black
- Perform a rotation at grandparent p



```

Del_FixUp(Ref root, Ref x) {
    while (x->color == BLACK && x != root)
        if (x == x->parent->left)
            del_leftAdjust(root, x);
        else
            del_rightAdjust(root, x);
    x->color = BLACK;
}

del_leftAdjust(Ref & root, Ref & x) {
    w = x->parent->right;
    if (w->color == RED) {
        w->color = BLACK;
        x->parent->color = RED;
        leftRotate(root, x->parent);
        w = x->parent->right;
    }
}

```

```

if (w->right->color == w->left->color == BLACK) {
    w->color = RED;
    x = x->parent;
}
else {
    if (w->right->color == BLACK) {
        w->left->color = BLACK;
        w->color = RED;
        rightRotate(root, w);
        w = x->parent->right;
    }
    w->color = x->parent->color;
    x->parent->color = w->right->color = BLACK;
    leftRotate(root, x->parent);
    x = root;
}
}


```

B-Trees

A B-tree of order t is either empty, or has the following properties:

- Every node contains at most $2t$ keys
- All nodes, except the root, contain at least t keys
- Every node is either a leaf or it has $m + 1$ descendants, where m is its number of keys
- All leaves are on the same level

Example



A Node (or Page) in B-Trees

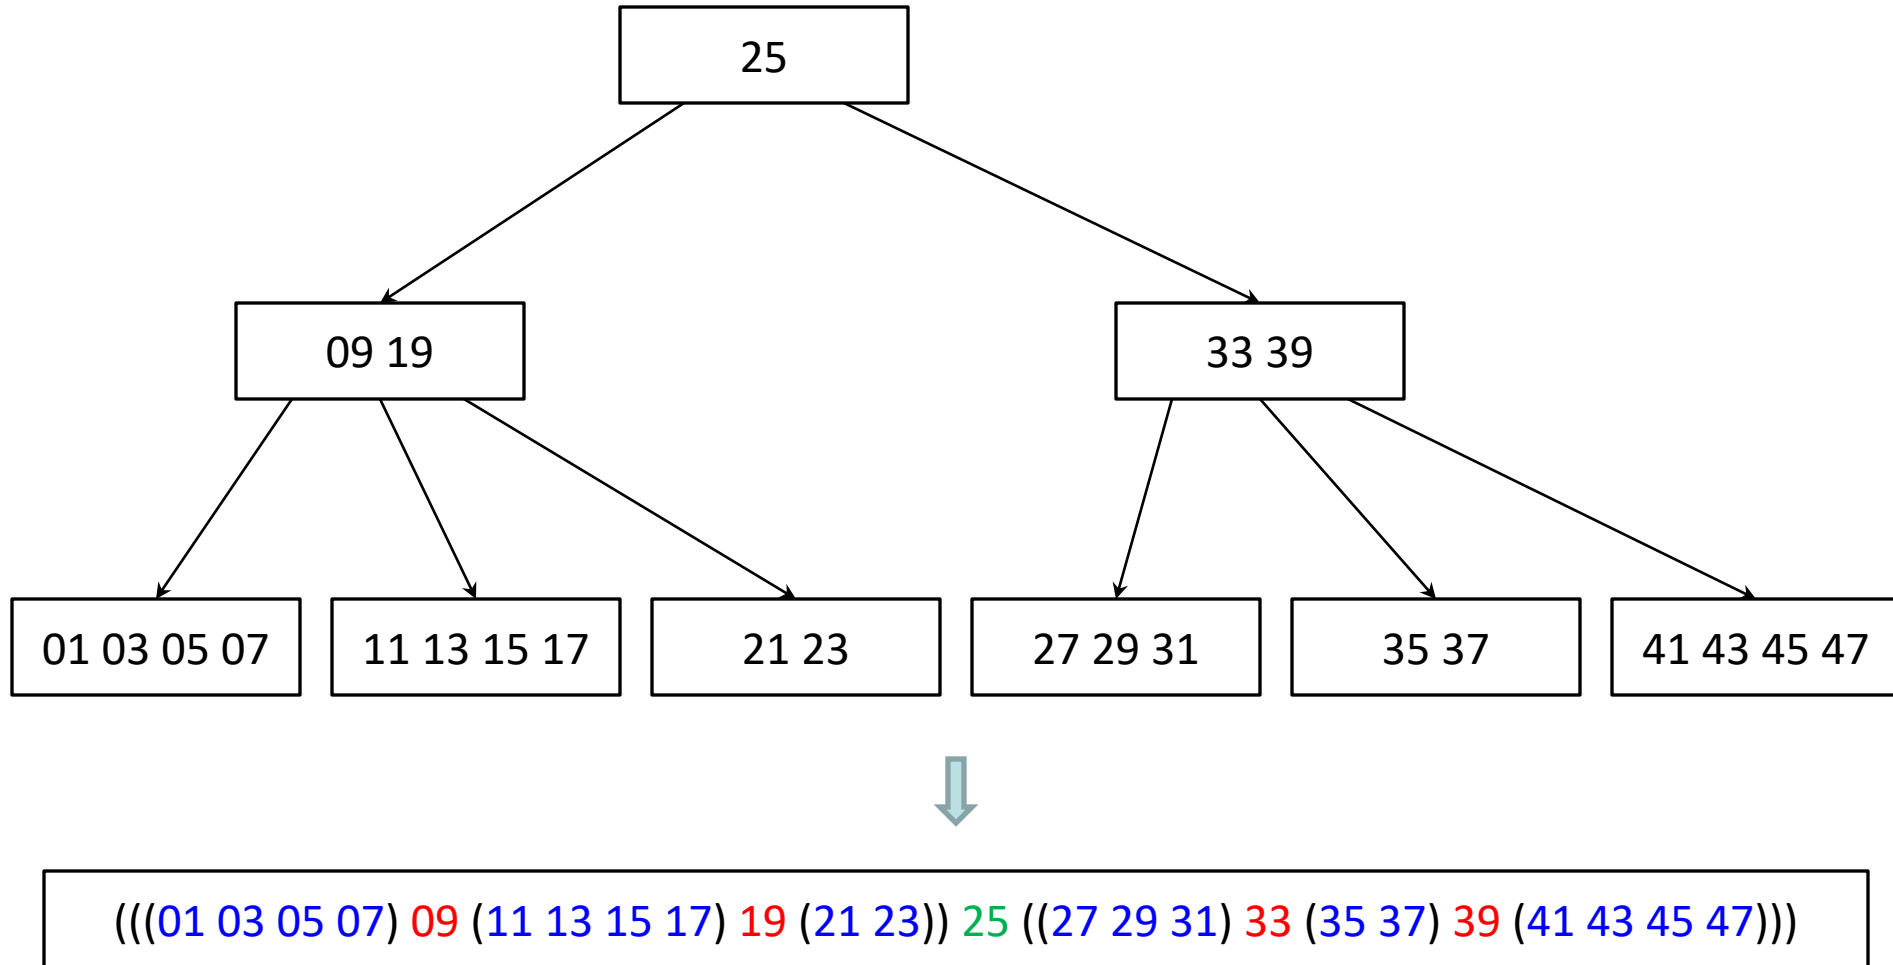
The form of a node is as follows:

p_0	k_1	p_1	k_2	p_2	\dots	k_m	p_m
-------	-------	-------	-------	-------	---------	-------	-------

where

- $k_1 < k_2 < \dots < k_m$
- p_i is a pointer to a descendant
 - If it's a leaf: $p_i = \text{NULL}, \forall i \in [0, m]$
- All keys in the node to which p_i points are greater than k_i and less than k_{i+1}

Example



Search Operation

- Let k be the search key. Assume that the node being considered contains m keys: k_1, k_2, \dots, k_m
- The search must start at the root of tree
 - If m is sufficiently large, one may use binary search; otherwise, a sequential search will do
- If the search is unsuccessful:
 - $k < k_1$: Search the node pointed to by p_0
 - $k > k_m$: Search the node pointed to by p_m
 - $k_i < k < k_{i+1}$: Search the node pointed to by p_i
- If the designated pointer is a null pointer: Stop!!!

Insertion Operation

- Assume that the key to be inserted is new
 - ➡ The search process terminates at a leaf
- The new key is inserted into the leaf if there is room
- If the leaf is full
 - *Insert the new key into the leaf*
 - Split the leaf into two nodes
 - Move the median key to the parent node
- The splitting can propagate upward up to the root, causing the tree to increase in height
 - This is the only way that a B-tree may increase its height

Deletion Operation

- Assume that the key to be deleted, say k , is in the tree
 - ➡ The search process terminates at the node containing k
- There are two different circumstances:
 - It's a leaf: The removal algorithm is plain and simple
 - Otherwise: The key must be replaced by its predecessor or successor, which happen to be on leaves and can easily be deleted
- ➡ In either cases, the key that actually to be deleted is always on a leaf

Deletion Operation

- If the leaf contains more than t keys
 - Delete k and no further action is required
 - If the leaf contains only t keys
 - If one of the adjacent siblings has more than t keys: Move one key from that sibling to the parent and one key from the parent to the leaf, and then delete k
 - Otherwise: Combine one of the adjacent siblings with the leaf and the median key from the parent, and then delete k
- ➡ This process may propagate all the way up to the root which could result in reducing the height of the B-tree

B-Trees: Summary

- In practice, B-trees are designed to store and manage a large data on secondary storage devices
- A node of a B-tree normally corresponds to a *disk page*
 - For a typical disk, a page might be 2^{11} to 2^{14} bytes in length
- The time needed to access a disk page is typically $\sim 10^5$ larger than the time needed to compare keys in RAM
 - The number of disk accesses (or the height of B-trees) is the principal indicator of the efficiency of this data structure